

LLM Platform Sample Document

This is a sample PDF used to test the RAG ingestion pipeline. It contains enough text for extraction, chunking, embeddings, and retrieval.

Large language models cannot efficiently search large documents directly. A Retrieval-Augmented Generation (RAG) pipeline extracts text, splits it into chunks, generates embeddings for each chunk, stores them in a vector database such as pgvector, and retrieves the most relevant chunks at query time.

This document should be successfully uploaded, parsed, chunked, and stored by the FastAPI RAG service.